

# MCB 5430 final assignment

Bianca Mocanu  
December 14, 2017

## Contents

<b>1</b>	<b>RNA-seq processing</b>	<b>1</b>
1.1	Pipeline for RNAseq data and generating bedgraph files . . . . .	1
1.2	Comparison of replicates . . . . .	1
<b>2</b>	<b>Differential Gene Expression Analysis</b>	<b>3</b>
2.1	EdgeR quantification of differences between datasets . . . . .	3
2.2	PCA plot . . . . .	7
2.3	MA plot . . . . .	8
2.4	Up and down-regulated genes . . . . .	10
2.5	Gene Ontology analysis using GSeq . . . . .	12
<b>3</b>	<b>Comparison of regulated genes with ChIP-seq data</b>	<b>16</b>
3.1	Bedtools for distances between ChIP peaks and nearest gene TSSs . . . . .	16
3.2	CDF plot . . . . .	18
3.3	Box plots . . . . .	19
3.4	The nature of the mechanism of gene activation / repression by ER . . . . .	20
3.5	UCSC Browser shot . . . . .	21
<b>4</b>	<b>Bonus question - cat hairball DNA motif</b>	<b>21</b>

# 1 RNA-seq processing

## 1.1 Pipeline for RNAseq data and generating bedgraph files

The RNA-seq processing script can be also found [here](#). It has the following steps:

- fastqc analysis of the initial fastq files provided on the server
- quality trimming of the fastq files and fastqc analysis of the output
- alignment using HISAT2
- bedgraph files generation for UCSC (with header) and with no header
- stringtie read quantification
- prepDE.py for edgeR compatible tables
- logfile

## 1.2 Comparison of replicates

First, the csv document from prepDE is imported as a matrix in R. A quick look at the matrix shows that to access the values for the E2 replicates, we need the first two columns, and to access the untreated replicates, we need columns 3 and 4. The log2 scatterplots can then be obtained by using the plot function with X and Y values as the log2 of the 1st and 2nd column values, respectively, for E2. Likewise, the 3rd and 4th columns are used to plot the same thing for untreated:

```
> #Starting prepDE counts files are in this directory:
>
> directory=~/.5430_RNAseq/prepDE/"
> setwd(directory)
> # Importing the csv file gene_count_matrix
>
> countData <- as.matrix(read.csv("gene_count_matrix.csv", row.names="gene_id"))
> head(countData)
```

	E2_rep1	E2_rep2	untr_rep1	untr_rep2
ENSG00000116032.5	8	13	11	15
ENSG00000137288.5	574	395	622	426
ENSG00000167578.12	394	326	187	288
ENSG00000102081.9	1338	1341	1794	1713
ENSG00000167531.2	0	1	0	1
ENSG00000103227.14	635	452	1150	514

```
> dfcountData <- data.frame(countData)
> plot(log2(dfcountData[,1]), log2(dfcountData[,2]), pch=20, col='gray60',
+       cex=0.2, main="Log2 scatterplot of the E2 replicates",
```

```


+       xlab="Log2(E2 rep 1 counts)",
+       ylab="Log2(E2 rep 2 counts)")
> # When adding the fit line and finding the correlation, there is a problem
> # for genes whose counts are zero. To fix this, I added 0.1 to the counts
> # number for those commands (does not alter the countData matrix).
>
> model1 <- lm(log2(dfcountData[,1]+0.1) ~ log2(dfcountData[,2]+0.1))
> abline(model1, col="red")
> model1_cor <- cor(log2(dfcountData[,1]+0.1), log2(dfcountData[,2]+0.1))
> model1_cor

[1] 0.9703333

> text(14, 1, paste("R squared =", round(model1_cor, digits=4)))

```

**Log2 scatterplot of the E2 replicates**



```

> # Same for the untreated replicates 1 and 2:
>
> plot(log2(dfcountData[,3]), log2(dfcountData[,4]), pch=20, col='gray60',
+       cex=0.2, main="Log2 scatterplot of the untreated replicates",
+       xlab="Log2(untr_rep 1 counts)",

```

```


+       ylab="Log2(untr_rep 2 counts)")
> model2 <- lm(log2(dfcountData[,3]+0.1) ~ log2(dfcountData[,4]+0.1))
> abline(model2, col="red")
> model2_cor <- cor(log2(dfcountData[,3]+0.1), log2(dfcountData[,4]+0.1))
> model2_cor

[1] 0.9642606

> text(14, 1, paste("R squared =", round(model2_cor, digits=4)))

```

**Log2 scatterplot of the untreated replicates**



## 2 Differential Gene Expression Analysis

### 2.1 EdgeR quantification of differences between datasets

For the EdgeR step, a multi dimensional DGEList object must be created. The DGEList function that creates the object requires as input the matrix and a "group" object that, in this case, groups the datasets based on whether they are the untreated controls or the treated samples.

In order to identify differentially expressed genes as accurately as possible, transcripts with very low counts are being excluded, as they are more error-prone. Here, only the genes

having more than one read per million in 3 out of 4 samples are being kept. Next, the data is normalized and the common and tagwise dispersions and variance values are calculated (tagwise, as required by the final exam questions). According to the results below, overall each data point can vary up to about 25 percent between replicates. The tagwise variance values are for each individual data point (here, 12719 of them).

The DE genes are called using the tagwise dispersion model, and the resulting table is filtered to keep entries with FDR values below 0.05. The remaining data from the remaining 394 genes can then later be highlighted in MA and MDS plots (see next sections).

```
> library(edgeR)
> library(car)
> # Get the name of the columns for the samples
>
> sampleNames <- colnames(countData)
> # Take a look at how countData object is displayed
>
> head(countData)
```

	E2_rep1	E2_rep2	untr_rep1	untr_rep2
ENSG00000116032.5	8	13	11	15
ENSG00000137288.5	574	395	622	426
ENSG00000167578.12	394	326	187	288
ENSG00000102081.9	1338	1341	1794	1713
ENSG00000167531.2	0	1	0	1
ENSG00000103227.14	635	452	1150	514

```
> # Creating groups (all samples that are untreated, all samples that are
> # treated) to be used for DGEList object later. The group syntax is such
> # that you name the label for the first group and then how many columns
> # are grouped, then same for 2nd group. Must be in that order that you
> # see in head(countData). Here, first two columns are the treated ones,
> # next 2 columns are the untreated ones.
>
> groups=c(rep("treated",2), rep('untreated',2))
> # DGEList function creates the DGEList object for the edgeR. We call it
> # cds here
>
> cds <- DGEList(countData, group=groups)
> names(cds) # this displays the components of the cds object
```

```
[1] "counts" "samples"
```

```
> # the cds object right now has the dataset with the counts and also some
> # information about the groups of samples that you have
> head(cds$counts)
```

	E2_rep1	E2_rep2	untr_rep1	untr_rep2
ENSG00000116032.5	8	13	11	15

```

ENSG00000137288.5      574      395      622      426
ENSG00000167578.12     394      326      187      288
ENSG00000102081.9     1338     1341     1794     1713
ENSG00000167531.2        0        1        0        1
ENSG00000103227.14     635     452     1150     514

> head(cds$samples)

      group lib.size norm.factors
E2_rep1   treated 20671047         1
E2_rep2   treated 17738490         1
untr_rep1 untreated 17051619         1
untr_rep2 untreated 15544963         1

> # Filtering poorly expressed genes
> cds <- cds[rowSums(1e+06 * cds$counts/expandAsMatrix(cds$samples$lib.size,
+                                                     dim(cds)) > 1) >=3, ]
> dim(cds)

[1] 12719      4

> cds <- calcNormFactors(cds)
> cds$samples

      group lib.size norm.factors
E2_rep1   treated 20671047  0.9819417
E2_rep2   treated 17738490  0.9814447
untr_rep1 untreated 17051619  1.0163386
untr_rep2 untreated 15544963  1.0209631

> #Effective library sizes after normalization
>
> cds$samples$lib.size * cds$samples$norm.factors

[1] 20297762 17409348 17330218 15870834

> cds <- estimateCommonDisp(cds)
> cds <- estimateTagwiseDisp(cds, prior.df = 10)
> names(cds) # now cds has much more info than before

[1] "counts"          "samples"          "common.dispersion"
[4] "pseudo.counts"   "pseudo.lib.size"  "AveLogCPM"
[7] "prior.df"        "prior.n"          "tagwise.dispersion"
[10] "span"

> # View common dispersion
> cds$common.dispersion

[1] 0.06696863

```

```

> # View tagwise dispersion values for several datapoints
> head(cds$tagwise.dispersion)

[1] 0.05933810 0.08756400 0.02700631 0.06424450 0.11769392 0.07335725

> tag_variance <- sqrt(as.numeric(cds$tagwise.dispersion))
> # View tagwise variance values for several datapoints
> head(tag_variance)

[1] 0.2435941 0.2959121 0.1643360 0.2534650 0.3430655 0.2708454

> common_variance <- sqrt(as.numeric(cds$common.dispersion))
> # View overall variance
> common_variance

[1] 0.258783

> # Using tagwise dispersion model to call significantly changed genes
> de.tgw <- exactTest(cds, dispersion = cds$tagwise.dispersion,
+                     pair = c("untreated", "treated"))
> head(de.tgw$table)

              logFC  logCPM    PValue
ENSG00000137288.5 -0.2994285 4.834848 0.40563805
ENSG00000167578.12 0.3988901 4.077626 0.36762094
ENSG00000102081.9 -0.5650186 6.470863 0.01919016
ENSG00000103227.14 -0.7864615 5.289378 0.03545942
ENSG00000188385.7 -0.1026021 2.148375 0.85714033
ENSG00000078237.4 -0.4883074 4.198229 0.22446361

> resultsTbl.tgw <- topTags(de.tgw, n=nrow(de.tgw$table))$table
> head(resultsTbl.tgw)

              logFC  logCPM    PValue    FDR
ENSG00000145439.7 -4.164982 6.645059 9.548435e-44 1.214465e-39
ENSG00000164626.8  4.733557 5.137123 9.437399e-35 6.001714e-31
ENSG00000136997.10 3.061029 7.457766 1.212849e-32 5.142077e-29
ENSG00000146243.9  4.714053 5.548333 8.016414e-31 2.549019e-27
ENSG00000113739.6  2.734506 8.182189 5.915024e-30 1.504664e-26
ENSG00000164128.2  3.805286 5.407338 1.262955e-28 2.677254e-25

> #Creating new table with significantly changed genes - here entries with
> # FDR values above 0.05 are being discarded
> de.genes.tbl.tgw <- resultsTbl.tgw[ resultsTbl.tgw$FDR <= 0.05, ]
> # This displays how many genes have FDR <= 0.05
> dim(de.genes.tbl.tgw)

[1] 394  4

```

## 2.2 PCA plot

The PCA plot, just like the MDS plot, can display the relationship between samples. They differ slightly in the way they define the "distance" between samples. As per the final exam request, this subsection shows the steps for a PCA plot. Despite previous analysis being done using EdgeR, DESeq2 package can build an object that can be then modified and used as an argument of DESeq2's plotPCA function, as seen below. At this point, the R environment already has the colData matrix from previous commands. The res object is somewhat equivalent to EdgeR's de.genes table that has been filtered for the FDR 0.05 upper threshold.

The PCA plot here considers the two most informative dimensions to show how the samples differ from each other. For example, on the X axis (1st PC), the untreated samples are very close to each other, which can be interpreted as a high degree of similarity, while the treated ones are a bit farther from each other.

```
> library(DESeq2)
> condition <- c("treated", "treated", "untreated", "untreated")
> colData <- data.frame(condition)
> row.names(colData) <- sampleNames
> colData

      condition
E2_rep1   treated
E2_rep2   treated
untr_rep1 untreated
untr_rep2 untreated

> all(rownames(sampleNames) %in% colnames(countData))


[1] TRUE

> countData <- countData[, sampleNames]
> all(rownames(colData) == colnames(countData))

[1] TRUE

> dds <- DESeqDataSetFromMatrix(countData = countData, colData = colData,
+                               design = ~ condition)
> dds <- DESeq(dds)
> res <- results(dds, alpha=0.05)
> rld <- rlogTransformation(dds, blind=TRUE)
> plotPCA(rld, intgroup=c('condition'))
```





## 2.3 MA plot

For the MA plot, two approaches are possible: either using a standard plot function or edgeR's integrated `plotSmear` function.


For the simple plot function, the log of the Fold Change can be plotted against the total counts per million from the initial output dataset of edgeR (`resultsTbl.tgw`). Then, we can overlay this plot with the same log of the Fold Change against the total counts per million from the filtered dataset (for the  $FDR \leq 0.05$ ) and display these data points e.g. in red, to show which of these genes are significantly changed. For guidance, horizontal lines for  $y = 1$  and  $-1$  can be included which, on the  $\log_2$  Y axis, indicate where a fold change of 2 and  $1/2$ , respectively, would be.

One observation here is that some genes can be confidently called as being differentially expressed even though their change is less than doubled, when upregulated, or less than halved, when downregulated. This is possible because the total counts per million are high enough so that the result is more confident than for low expressed genes.

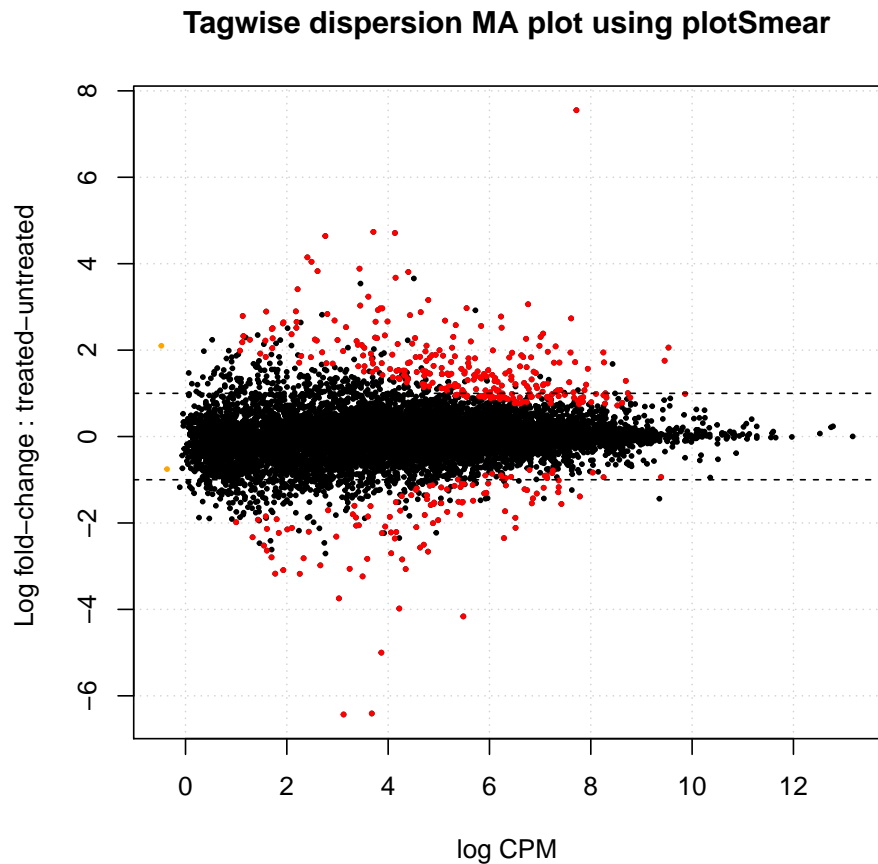
```
> # The MA plot using plot function
>
```

```
> plot(resultsTbl.tgw$logCPM, resultsTbl.tgw$logFC, pch=20, col='grey60',
+       cex=0.5, main = "Tagwise dispersion MA plot using plot() and points()",
+       xlab="log2(total CPM)",
+       ylab="log2(Fold Change)")
> # adding horizontal lines for y = -1 and 1
> abline(h=-1, lty=5)
> abline(h=1, lty=5)
> points(de.genes.tbl.tgw$logCPM, de.genes.tbl.tgw$logFC, pch=20, col='red',
+        cex=0.75)
```

**Tagwise dispersion MA plot using plot() and points()**



```
> # The MA plot using EdgeR's plotSmear function
>
> de.genes.tgw <- rownames(resultsTbl.tgw)[resultsTbl.tgw$FDR <= 0.05]
> plotSmear(cds, de.tags=de.genes.tgw,
+           main="Tagwise dispersion MA plot using plotSmear",
+           pair=c("untreated", "treated"), cex=0.35, xlab="log CPM",
+           ylab="Log fold-change")
> abline(h=-1, lty=2)
> abline(h=1, lty=2)
```



## 2.4 Up and down-regulated genes

To obtain the up and downregulated gene name lists, the starting point is the output file of EdgeR, `resultsTbl.tgw`.

The upregulated and downregulated genes are defined as entries with the FDR value  $\leq 0.05$  which have logFC values above zero or below zero, respectively. These entries, however, have row names that correspond to the transcript variants of the genes, as indicated by the number after the separating ".". To reach to the gene name alone, a splitting function must be applied to all rows of the Up and Downregulated genes list.

```
> library(rtracklayer)
> write.table(resultsTbl.tgw , file="all_genes_tgw.txt", col.names =T,
+             row.names=T)
> resultsTbl.tgw <- read.table("all_genes_tgw.txt")
> #filter based on FDR
> de.genes.tgw <- resultsTbl.tgw[ resultsTbl.tgw$FDR <= 0.05,]
> dim(de.genes.tgw)
```

```
[1] 394  4
```

```

> E2_up <- de.genes.tgw [de.genes.tgw $logFC >0,]
> E2_down <- de.genes.tgw [de.genes.tgw $logFC <0,]
> dim(E2_up)

[1] 280    4

> dim(E2_down)

[1] 114    4

> Genes.all <- rownames(resultsTbl.tgw)
> E2_up_names <- rownames(E2_up)
> E2_down_names <- rownames(E2_down)
> # The above gene names still have the transcript variant attached after
> # the "." separator. In order to obtain
> # The gene name, they need to be split:
>
> Genes.all_split <- sapply(Genes.all, function(x)
+   unlist(strsplit(x, split=".[.]"))[1])
> E2up_split <-sapply(E2_up_names, function(x)
+   unlist(strsplit(x, split=".[.]"))[1])
> E2down_split <-sapply(E2_down_names, function(x)
+   unlist(strsplit(x, split=".[.]"))[1])
> write.table(as.data.frame(Genes.all_split), file="Genes_all_split.txt",
+   row.names=TRUE, col.names=TRUE)
> write.table(as.data.frame(E2up_split), file="E2_upregulated_split.txt",
+   row.names=TRUE, col.names=TRUE)
> write.table(as.data.frame(E2down_split), file="E2_downregulated_split.txt",
+   row.names=TRUE, col.names=TRUE)
> head(Genes.all_split)

ENSG000000145439.7  ENSG000000164626.8  ENSG000000136997.10  ENSG000000146243.9
"ENSG000000145439"  "ENSG000000164626"  "ENSG000000136997"  "ENSG000000146243"
ENSG000000113739.6  ENSG000000164128.2
"ENSG000000113739"  "ENSG000000164128"

```

Following the sapply function, however, the lists and text files still contain the initial gene name (with "."), as well as the gene name. An easy fix is using the following commands in the terminal to cut this space delimited text, remove the "" surrounding gene names and retrieve only the gene names, and eventually remove the precursor text file.

```

cat Genes_all_split.txt | cut -d " " -f2 | sed 's/[^a-z A-Z 0-9]//g'
> Gene_names_all.txt

# rm Genes_all_split.txt

cat E2_upregulated_split.txt | cut -d " " -f2 | sed 's/[^a-z A-Z 0-9]//g'
> Upregulated_names.txt

```

```
# rm E2_upregulated_split.txt

cat E2_downregulated_split.txt | cut -d " " -f2 | sed 's/[^a-z A-Z 0-9]
//g' > Downregulated_names.txt

# rm E2_downregulated_split.txt
```

## 2.5 Gene Ontology analysis using GOseq


```
> library(goseq)
> library(org.Hs.eg.db)
> up.vector=as.integer(Genes.all_split%in%E2up_split)
> down.vector=as.integer(Genes.all_split%in%E2down_split)
> head(up.vector)

[1] 0 1 1 1 1 1

> # Add names from all genes to vector
>
> names(up.vector)<-Genes.all_split
> names(down.vector)<-Genes.all_split
> head(up.vector)

ENSG000000145439 ENSG000000164626 ENSG000000136997 ENSG000000146243 ENSG000000113739
                0                1                1                1                1
ENSG000000164128
                1

> pwf_up=nullp(up.vector,"hg19","ensGene")
>
```



```
> pwf_down=nullp(down.vector,"hg19","ensGene")
> head(pwf_up)
```

	DEgenes	bias.data	pwf
ENSG00000145439	0	2128	0.01996940
ENSG00000164626	1	3787	0.02889231
ENSG00000136997	1	1346	0.01456005
ENSG00000146243	1	1161	0.01314117
ENSG00000113739	1	5347	0.02983885
ENSG00000164128	1	2762	0.02423019

```
> GO.wall_up=goseq(pwf_up,"hg19","ensGene", method = "Wallenius")
> GO.wall_down=goseq(pwf_down,"hg19","ensGene", method = "Wallenius")
> head(GO.wall_up)
```

	category	over_represented_pvalue	under_represented_pvalue	numDEInCat
4484	G0:0010604	3.312885e-07	1.0000000	80
4299	G0:0009893	1.809705e-06	0.9999994	82
12191	G0:0048534	1.920803e-06	0.9999994	29
17418	G0:1902107	4.554821e-06	0.9999994	10

1026	GO:0002520	6.967464e-06	0.9999976	29
18068	GO:1903708	8.658716e-06	0.9999987	11
	numInCat			term ontology
4484	2044	positive regulation of macromolecule metabolic process		BP
4299	2183	positive regulation of metabolic process		BP
12191	461	hematopoietic or lymphoid organ development		BP
17418	70	positive regulation of leukocyte differentiation		BP
1026	492	immune system development		BP
18068	91	positive regulation of hemopoiesis		BP

```
> head(GO.wall_down)
```

		category	over_represented_pvalue	under_represented_pvalue	numDEInCat
3367	GO:0007179		2.198142e-05	0.9999975	8
15201	GO:0071560		8.446198e-05	0.9999883	8
15200	GO:0071559		8.897599e-05	0.9999876	8
10067	GO:0042704		1.095039e-04	1.0000000	2
3366	GO:0007178		1.097169e-04	0.9999818	9
6784	GO:0030278		1.266944e-04	0.9999844	7
	numInCat				
3367	128				
15201	155				
15200	156				
10067	2				
3366	204				
6784	123				
					term
3367		transforming growth factor beta receptor signaling pathway			
15201		cellular response to transforming growth factor beta stimulus			
15200		response to transforming growth factor beta			
10067		uterine wall breakdown			
3366		transmembrane receptor protein serine/threonine kinase signaling pathway			
6784		regulation of ossification			
	ontology				
3367	BP				
15201	BP				
15200	BP				
10067	BP				
3366	BP				
6784	BP				

```
> enriched.GO.BP_up=GO.wall_up$category[
+   p.adjust(GO.wall_up$over_represented_pvalue, method="BH")<0.05]
> #The FDR of 0.01 would not give many results here, so I had to use 0.05 and worse.
>
> enriched.GO.BP_down=GO.wall_down$category[
+   p.adjust(GO.wall_down$over_represented_pvalue, method="BH")<0.1]
> head(enriched.GO.BP_up)
```

```
[1] "GO:0010604" "GO:0009893" "GO:0048534" "GO:1902107" "GO:0002520"
[6] "GO:1903708"
```

```
> head(enriched.GO.BP_down)
```

```
character(0)
```

```
> library(GO.db)
> for(go in enriched.GO.BP_up[1:5]){
+ print(GOTERM[[go]])
+ cat("-----\n")
+ }
```

```
GOID: GO:0010604
```

```
Term: positive regulation of macromolecule metabolic process
```

```
Ontology: BP
```

```
Definition: Any process that increases the frequency, rate or extent of
the chemical reactions and pathways involving macromolecules, any
molecule of high relative molecular mass, the structure of which
essentially comprises the multiple repetition of units derived,
actually or conceptually, from molecules of low relative molecular
mass.
```

```
-----
GOID: GO:0009893
```

```
Term: positive regulation of metabolic process
```

```
Ontology: BP
```

```
Definition: Any process that activates or increases the frequency, rate
or extent of the chemical reactions and pathways within a cell or
an organism.
```

```
Synonym: activation of metabolic process
```

```
Synonym: positive regulation of metabolism
```

```
Synonym: stimulation of metabolic process
```

```
Synonym: up regulation of metabolic process
```

```
Synonym: up-regulation of metabolic process
```

```
Synonym: upregulation of metabolic process
-----
```

```
GOID: GO:0048534
```

```
Term: hematopoietic or lymphoid organ development
```

```
Ontology: BP
```

```
Definition: The process whose specific outcome is the progression of
any organ involved in hematopoiesis (also known as hemopoiesis) or
lymphoid cell activation over time, from its formation to the
mature structure. Such development includes differentiation of
resident cell types (stromal cells) and of migratory cell types
dependent on the unique microenvironment afforded by the organ for
their proper differentiation.
```

```
Synonym: haematopoietic or lymphoid organ development
```

```
Synonym: haemopoietic or lymphoid organ development
```



Synonym: hemopoietic or lymphoid organ development

GOID: GO:1902107

Term: positive regulation of leukocyte differentiation

Ontology: BP

Definition: Any process that activates or increases the frequency, rate or extent of leukocyte differentiation.

Synonym: activation of immune cell differentiation

Synonym: activation of leucocyte differentiation

Synonym: activation of leukocyte differentiation

Synonym: positive regulation of immune cell differentiation

Synonym: positive regulation of leucocyte differentiation

Synonym: up regulation of immune cell differentiation

Synonym: up regulation of leucocyte differentiation

Synonym: up regulation of leukocyte differentiation

Synonym: up-regulation of immune cell differentiation

Synonym: up-regulation of leucocyte differentiation

Synonym: up-regulation of leukocyte differentiation

Synonym: upregulation of immune cell differentiation

Synonym: upregulation of leucocyte differentiation

Synonym: upregulation of leukocyte differentiation

GOID: GO:0002520

Term: immune system development

Ontology: BP

Definition: The process whose specific outcome is the progression of an organismal system whose objective is to provide calibrated responses by an organism to a potential internal or invasive threat, over time, from its formation to the mature structure. A system is a regularly interacting or interdependent group of organs or tissues that work together to carry out a given biological process.

### 3 Comparison of regulated genes with ChIP-seq data

#### 3.1 Bedtools for distances between ChIP peaks and nearest gene TSSs

To calculate these distances, the following steps need to be done and considered:

- Use bedtools closest to report the distances between the summits (-a) and gene TSSs (-b) with respect to the TSS positions (-D b).
- To be able to plot this information later, only one TSS must be printed in the event of a tie (-t first)
- To get the TSSs, start with a bed file with all the genes and reformat it (e.g. using awk) so that the transcription start site is either the 2nd or 3rd column in the bed file,

depending on the strand where the gene is.

- I don't think there is a need for further processing like flanking the TSSs by e.g. 50 nucleotides each way, since bedtools closest reports only one distance from a summit to the closest TSS
- Since further analysis asks for comparison with the up, down, and non-regulated genes, the gene names must be the Ensemble IDs "ENSG[...]" in order to be compatible with the text files with all the genes retrieved in Part 2, section 4.

```
awk '{if($6=="+")
> printf ("%s\t%s\t%s\t%s\t%s\t%s\n", $1, $2, $2, $4, $5, $6);
> else
> printf ("%s\t%s\t%s\t%s\t%s\t%s\n", $1, $3, $3, $4, $5, $6)
> }' hg19_gencode_ENSG_geneID.bed > hg19_gencode_TSS_ENSG_geneID.bed
```

The file with the TSS then needs to be sorted. The sorted file can be used with bedtools closest with the -D b and -t first options:

```
bedtools sort -i hg19_gencode_TSS_ENSG_geneID.bed >
hg19_gencode_TSS_ENSG_geneID_sorted.bed

rm hg19_gencode_TSS_ENSG_geneID.bed

bedtools closest -D b -t first -a /tempdata3/MCB5430/final_data/ER_peaks
/ER_E2_summits.bed -b hg19_gencode_TSS_ENSG_geneID_sorted.bed >
ER_E2_closest_all.bed
```

```
> #loading the BED file into R
> E2_all_closest <- read.table("ER_E2_closest_all.bed",
+ header=FALSE, sep="\t", stringsAsFactors = FALSE)
> head(E2_all_closest)
```

	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10
1	chr1	856576	856577	MACS_peak_1	7	chr1	860259	860259	ENSG00000187634	591
2	chr1	1009194	1009195	MACS_peak_2	125	chr1	1009687	1009687	ENSG00000237330	592
3	chr1	1014951	1014952	MACS_peak_3	33	chr1	1009687	1009687	ENSG00000237330	592
4	chr1	1316033	1316034	MACS_peak_4	12	chr1	1310875	1310875	ENSG00000175756	74
5	chr1	1368560	1368561	MACS_peak_5	25	chr1	1370240	1370240	ENSG00000179403	595
6	chr1	2379468	2379469	MACS_peak_6	10	chr1	2357418	2357418	ENSG00000149527	75
	V11	V12								
1	+	-3682								
2	-	492								
3	-	-5264								
4	-	-5158								
5	+	-1679								
6	+	22050								

```

> #loading the lists of names in R
> all_gene_names <- read.table("Gene_names_all.txt", header=TRUE,
+                             stringsAsFactors = FALSE)
> upreg_names <- read.table("Upregulated_names_sorted.txt",
+                             header=TRUE, stringsAsFactors = FALSE)
> downreg_names <- read.table("Downregulated_names_sorted.txt",
+                             header=TRUE, stringsAsFactors = FALSE)
> # To subset the E2_all_closest, I first selected the rows on the
> # condition that V9 (the ENSG name) is within the gene list names, then
> # picked those rows from E2_all_closest
>
> selected_rows <- (E2_all_closest$V9 %in% upreg_names$E2upsplit)
> Upreg_closest <- E2_all_closest[selected_rows,]
> #this is the dataframe with only upregulated genes
>
> selected_rows <- (E2_all_closest$V9 %in% downreg_names$E2downsplit)
> Downreg_closest <- E2_all_closest[selected_rows,]
> selected_rows <- (E2_all_closest$V9 %in% all_gene_names$Genesallsplit)
> Allgenes_closest <- E2_all_closest[selected_rows,]
> # To obtain the non-regulated table, I then removed the regulated genes
> # rows from the E2_all_closest dataset
>
> library(dplyr)
> Regulated_closest <- union(Upreg_closest, Downreg_closest)
> Nonregulated_closest <- setdiff(Allgenes_closest, Regulated_closest)
>

```

### 3.2 CDF plot

The cumulative distribution plots shown here allow to estimate how far from the TSS most of the ChIP peaks are. The X axis is for the distance between the TSS and the ChIP peaks and the y axis shows what proportion of peaks are within x basepairs from the TSS. The bedtools closest generated distances that have + and -, so in order to make meaningful plots, the absolute value of the reported distance is used here.

```

> up_closest_cdf <- ecdf(abs(Upreg_closest[,12]))
> down_closest_cdf <- ecdf(abs(Downreg_closest[,12]))
> non_closest_cdf <- ecdf(abs(Nonregulated_closest[,12]))
> plot(up_closest_cdf, xlim=c(0,50000), ylim=c(0.1, 0.9),
+      col="forestgreen", main = "CDF plot of ER E2 distance to TSS",
+      xlab="Distance (bps)", ylab = "Fraction", las=1)
> par(new=TRUE)
> plot(down_closest_cdf, xlim=c(0,50000), ylim=c(0.1, 0.9),
+      col="red", main = "CDF plot of ER E2 distance to TSS",
+      xlab="Distance (bps)", ylab = "Fraction", las=1)
> par(new=TRUE)
> plot(non_closest_cdf, xlim=c(0,50000), ylim=c(0.1, 0.9),

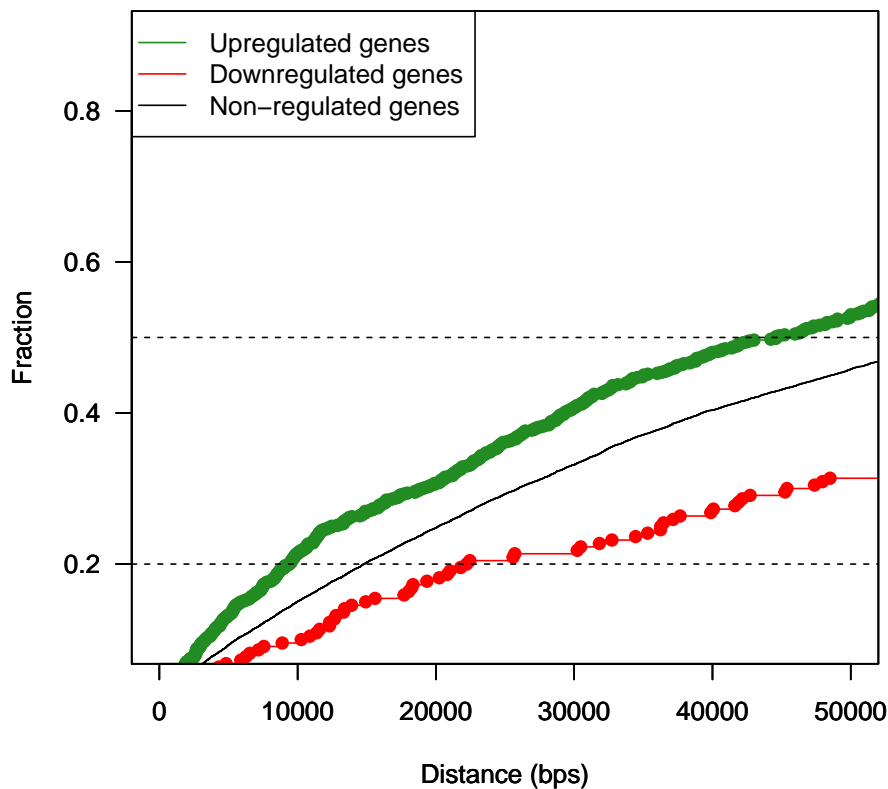
```

```

+      col="black", main = "CDF plot of ER E2 distance to TSS",
+      xlab="Distance (bps)", ylab = "Fraction", las=1)
> abline(h=0.5, lty=2)
> abline(h=0.2, lty=2)
> legend( x="topleft", legend=c("Upregulated genes",
+                               "Downregulated genes",
+                               "Non-regulated genes"),
+        col=c("forestgreen", "red", "black"), lwd=1, lty=c(1,1,1),
+        pch=c(NA,NA) )
>

```

**CDF plot of ER E2 distance to TSS**



### 3.3 Box plots

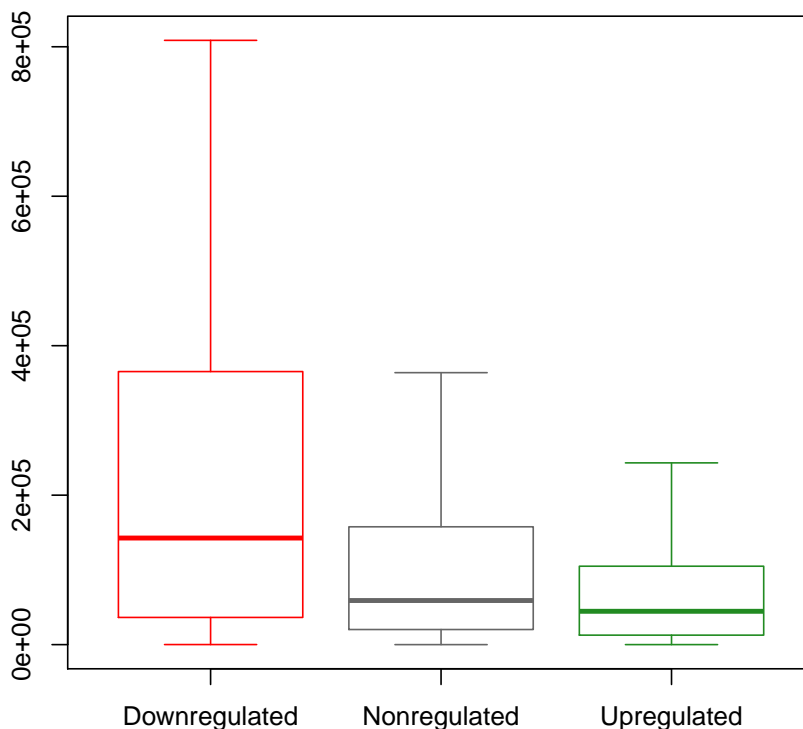
The starting point for boxplots are the `Upreg_closest`, `Downreg_closest` and `Nonregulated_closest`. These are currently dataframes that all have 12 variables and different number of rows (observations). Combining these separate datasets by columns doesn't work well (or serve the purpose), as they all have different number of rows. To be able to use combine them so that the boxplots are plotted together, I'd create a new column for each of these three, containing the type of effect ER has on the genes (upregulating, downregulating, or none), combine the tables with this new column, and plot the distances grouped by the effect.

```

> Upreg_closest$effect<-"Upregulated"
> Downreg_closest$effect<-"Downregulated"
> Nonregulated_closest$effect<-"Nonregulated"
> Genes_with_effects<- rbind(Upreg_closest, Downreg_closest, Nonregulated_closest)
> boxplot(abs(Genes_with_effects$V12) ~ Genes_with_effects$effect, outline=F,
+         border=c("red", "grey40", "forestgreen"),
+         main="Distribution of distances between ChIP peaks and gene TSS", lty=1)

```

### Distribution of distances between ChIP peaks and gene TSS



### 3.4 The nature of the mechanism of gene activation / repression by ER

According to the CDF plot, only approximately 20 percent of peaks are less than 10k basepairs away from the TSS for upregulated genes. For non-regulated genes and downregulated genes, it seems even worse.

The boxplots reveal the same thing, as the peaks for downregulated genes seem to be much farther away from gene TSS than for non-regulated and upregulated genes, as well as much more dispersed.

These results are expected for the estrogen receptor, given the previous ChIP seq analyses done for the midterm which showed that this receptor has a very strong preference for enhancers and binds very little to TSS. This is not consistent with what I would expect for a canonical transcription factor (such as ATF1). I would think that a transcription factor prefers areas

around gene TSS and binds directly to those, which would result in a very steep CDF line that would look asymptotic at a much higher fraction (much closer to 1), within 50000 basepairs. The boxplots would also look much more condensed.

### 3.5 UCSC Browser shot

To find an upregulated gene on Chromosome 12, I used the following commands:

```
cat Upregulated_names.txt | while read line
> do
> grep $line hg19_gencode_ENSG_geneID.bed >> upgenesssss.bed
> done
```

chr12	119616446	119632551	ENSG000000152137	1497	+
chr12	123468026	123594975	ENSG000000090975	190	-
chr12	31433517	31479050	ENSG000000139146	12	-
chr12	116396380	116715143	ENSG000000123066	184	-
chr12	76419226	76425384	ENSG000000139289	1168	-
chr12	64238072	64541613	ENSG000000196935	134	+
chr12	2904118	2913124	ENSG00000004478	607	+

The second gene in the list above is displayed in the attached browser shot. Its RefSeq name is PITPNM2.

## 4 Bonus question - cat hairball DNA motif

For the catball motif, I thought I could run blat or bwa on the GGGGGCCCCCAT sequence, get the coordinates, subtract or add 904,100 to the coordinates and then see what motifs can lie around. BWA and Blat seem to consider this sequence too short to allow that (they want at least 50 basepairs), so I uploaded the hairball dna text file as a fasta file in Benchling and I searched for the GGGGGCCCCCAT as a motif there - I found it to be at 3074588-3074599, which means the motif we are looking for is around position 2170488-2170499, as shown below:

